

Time-Inhomogeneous Multi-State Model for a Swine Flu Epidemic

October 17, 2012

1 Introduction

At the beginning of epidemics, such as H1N1 (“swine flu”) or severe acute respiratory syndrome (SARS), the major concern is the severity of the illness, and the large number of individuals that have the potential to be hospitalized after being infected. Accordingly, during the early stages of an epidemic we can assume that most cases are identified as ill individuals are encouraged to seek medical help and practitioners to send samples to the labs for testing. After the first few weeks of an epidemic, testing of sick individuals is often discouraged to ease the burden on public health labs.

In infectious disease modelling, we are often interested in estimating the case fatality ratio. That is, we are interested in the proportion of cases that eventually die from the disease. However, crude rates calculated throughout the epidemic often underestimate this quantity due to the large proportion of censoring present in the data. The reason for this being that many individuals have yet to reach the terminal state of their illness, either recovery or death. Additionally, these crude rates are often misleading due to shifts in case ascertainment which occur as the epidemic progresses because efforts from medical professionals become focussed on the more severe cases. If these biases are not accounted for, they result in underestimation of the case fatality ratio (?). In fact, ?, even omit performing analyses during the first two months of the SARS epidemic due to a lack of data.

We propose a Bayesian approach to this problem using time-inhomogeneous multi-state models. This approach takes into account expert opinion which allows us to get reasonable estimates at the beginning an epidemic when the data are sparse. The use of survival analysis is a logical approach because this is the standard way of dealing with the large amounts of censoring present in the data. We illustrate our approach through a simulation study with the aim of appropriately estimating the fatality rate, predicting the hospital load and predicting the number of unobserved cases over the course of the epidemic. We are also interested in forecasting the number of cases beyond the completion our simulated epidemic. The strength of this method is appropriately estimating these quantities at the beginning of the epidemic while the data are sparse and censored, rather than at the completion.

We also compare our methods to similar analyses...

This remainder of this paper proceeds as follows. Section 2 goes into detail describing our multi-state model and Section 3 discusses model inference. The results of our simulation

study are shown in Section 4. Finally, our conclusions are given in Section ??.

2 Time-Inhomogeneous Multi-State Models

2.1 Infection Types

We assumed that all individuals had one of three infection types: mild, serious or deadly. Individuals were classified as mild if they recovered without needing to be hospitalized. If an infection was severe enough that hospitalization was required before recovery, they were classified as having a serious infection type. Individuals with a deadly infection, as the name suggests, were those that died following hospitalization. For mild, serious and deadly cases, respectively, we modelled the disease progression as follows:

$$\text{Infection} \xrightarrow{f_I(t)} \text{Onset} \xrightarrow{f_{OM}(t)} \text{Medical} \xrightarrow{f_{MR}(t)} \text{Recovery} \quad (1)$$

$$\text{Infection} \xrightarrow{f_I(t)} \text{Onset} \xrightarrow{f_{OS}(t)} \text{Medical} \xrightarrow{f_{SH}(t)} \text{Hospitalization} \xrightarrow{f_{SR}(t)} \text{Recovery} \quad (2)$$

$$\text{Infection} \xrightarrow{f_I(t)} \text{Onset} \xrightarrow{f_{OD}(t)} \text{Medical} \xrightarrow{f_{HD}(t)} \text{Hospitalization} \xrightarrow{f_D(t)} \text{Death} \quad (3)$$

Regardless of an individual's infection type, progression from one stage to the next occurred according to a Weibull distribution with parameters that corresponded to the respective transition and infection type. Note that until an individual died or recovered their infection type was unknown. We also assumed that a small proportion of mild cases were lost to follow up. That is, they received medical consultation, and were assumed to recover since no further information was available.

In this analysis, we assumed that the infection type probabilities varied with respect to age which we generated using log-linear cubic splines. This will be discussed further in Section 2.3.

2.2 The Model

We used a zero-inflated rounded Weibull for the time-to-event distribution, denoted $f(t)$. We used a rounded distribution because the data only included the date the event took place. For example, if there was an event recorded as being on the first day of the epidemic, the event would actually have occurred at some point between day 0.5 and day 1.5. To account for the large number of administrative zeros recorded in the data, we used a zero-inflated distribution. That is, we let δ be the probability of transitioning immediately to the death stage for an administrative reason, such as arriving dead at the hospital. Furthermore, if we let Z denote the random variable for administrative zeros, we can assume Z follows a Bernoulli distribution with probability δ and hence,

$$\Pr(T = t | Z = 1) = 0 \text{ with probability } 1$$

For the case where the individual was not an administrative zero

$$\Pr(T = t|Z = 0) = \Pr[\max(t - 0.5, 0) < y < t + 0.5]$$

$$y \sim \text{Weibull} \left[\frac{\mu}{\Gamma(1 + 1/\nu)}, \nu \right]$$

where $\mu = E[T = t|Z = 0]$ is the average time to an event and ν is the Weibull shape parameter. We used beta priors for the zero-inflation parameters and gamma priors for the Weibull mean and shape parameters.

2.3 Infection Type Probabilities

As mentioned above in Section 2.1, we assumed the infection type probabilities varied with respect to age and generated them using log-linear cubic splines. For individual i we let $W_i \in \{M, S, D\}$ denote the infection type, where M , S and D represent mild, serious and deadly infections, respectively and X_i represents the individual's age. The fatality rate is the marginal probability $\Pr(W_i = D) = g_D(X_i)$ and the non-fatal hospitalization rate is the conditional probability $\Pr(W_i = S|W_i \neq D) = g_S(X_i)$ where g_D and g_S are log-linear cubic splines with fixed knots and random coefficients that vary smoothly with age. Note that this was done using the R package, DPpackage.




Figure 1: Death probabilities.

2.4 Discussion of Prior Distributions

3 Model Inference

Inference for this model was done through data augmentation and random walk Metropolis simulations for calculating event time parameters and predicting the number of unobserved cases. We also forecasted the number of cases beyond the end of the simulated epidemic.

3.1 Data Augmentation

In order to deal with the censoring, we performed data augmentation to impute the lifetimes as follows:

1. choose starting values for the probabilities
2. simulate $[W_i = D] \sim \text{Bernoulli}(p_i)$, where $p_i = \Pr(W_i = D | T_i > k_i; X_i)$, k_i is the number of days previous that individual i was hospitalized and X_i is the age of individual i
3. simulate $[g_D | W_1, \dots, W_n]$ to update the parameters

3.2 Event Time Parameters

After the data augmentation step, the likelihood was straightforward to calculate since the observations were assumed to be independent. Moreover, we could easily use random walk Metropolis to calculate the event time parameters. If we denote θ as the current step and ϕ as the proposed value, then $q(\theta, \phi)$ is the proposal distribution for moving from the current position θ to the proposed position ϕ . For this simulation the proposal distribution was absolute normal. The acceptance probability for the random walk was:

$$\alpha(\theta, \phi) = \min \left\{ 1, \frac{\pi(\phi)}{\pi(\theta)} \right\}$$

where $\pi(\cdot) \propto \ell(\cdot)p(\cdot)$. Using this notation $p(\cdot)$ is the prior distribution, $\ell(\cdot)$ is the likelihood and $\pi(\cdot)$ is the posterior distribution. We accepted the move if $u \leq \alpha$ where $u \sim \text{Unif}(0, 1)$ and rejected otherwise.

3.3 Number of Unobserved Cases

Again, we used random walk Metropolis to calculate the number of unobserved infections. We let y_j denote the number of observed cases on day j of a J day epidemic and we assumed that for each day of the simulation:

$$y_j \sim \text{Binomial}(N_j, p_j), j = 1, \dots, J$$

where N_j is the total number of cases (including both observed and unobserved) on the j^{th} day of the epidemic and p_j is the probability of observing a case on day j . A Poisson prior with mean λ_j is used for N_j , where λ_j is $[N_{j_{\inf}}\theta + \omega]\gamma$, $j = 1, \dots, J$. Using this notation, $N_{j_{\inf}}$ is the number of infective individuals on the j^{th} day of the epidemic, θ is the rate parameter, ω is the immigration parameter and γ is the probability of the case being mild, serious or deadly.

For this simulation, we accepted the proposed value with probability:

$$\alpha(\theta, \phi) = \min \left\{ 1, \frac{\pi(\phi)q(\phi, \theta)}{\pi(\theta)q(\theta, \phi)} \right\}$$

where $\pi(\cdot)$ is the posterior distribution and $q(\theta, \phi)$ is the Poisson proposal distribution for day j with mean $N_{j-1} + \pi_0$, where π_0 denotes the proposal offset or the mean increase in the number of new cases each day of the epidemic. Note that for this simulation $q(\theta, \phi)$ is the conditional distribution, given N_j is greater than $y_j - 0.5$. Hence, in order for $q(\phi, \theta)$ to be a valid probability density function we divided the conditional distribution by $P(Y_j > N_{j-1} + \pi_0)$, where Y_j are the individuals that have been infected, but have not had their medical consultation by day j . Similarly, for $q(\theta, \phi)$, we divided by $P(Y_j > N_j + \pi_0)$. We accepted the proposed value if $u \leq \alpha$ where $u \sim \text{Unif}(0, 1)$.

Once this was complete, we performed data augmentation to estimate the progression of the cases from stage to stage by infection type as described in the multi-state models, 1 - 3 (Section 2.1).

3.4 Forecasting the Number of Cases

The final goal of this paper was to predict the cases for a number of days beyond the completion of a simulated epidemic. For a J day epidemic, if we denote N_k , $k = J+1, \dots, K$, as the number of cases on the k^{th} day of the epidemic. Then,

$$N_k \sim \text{Poisson}([N_{k_{\inf}}\theta + \omega]\gamma), k = J+1, \dots, K$$

with $N_{k_{\inf}}$, θ , ω and γ as defined in Section 3.3. Note that these forecasts were done separately for each infection type. Once this was complete, we performed data augmentation, as discussed in Section 3.1.

4 Results of the Swine Flu Simulation Study

This section summarizes the results of a simulated 20 day epidemic and then forecasted the number of cases for the next 10 days beyond the end of the epidemic.

4.1 Fatality Rate